

Pham Van Tam

+84 941496803 | working@tamus.dev | github.com/twilightt1 | tamus.dev

SUMMARY

Computer Science student with a strong foundation in machine learning, deep learning, and applied NLP, and hands-on experience building LLM-based and retrieval-augmented systems. Comfortable working across the ML lifecycle, from data processing and model development to evaluation and deployment, with solid software engineering fundamentals including version control, testing, and clean code. Particularly interested in MLOps and LLM serving infrastructure, with a systems-oriented approach to real-world AI/ML engineering.

EDUCATION

Posts and Telecommunications Institute of Technology	Hanoi, Vietnam
<i>Engineer of Computer Science</i>	<i>Sep. 2024 – Present</i>
– Relevant Coursework: Machine Learning, Deep Learning, Probability and Statistics, Linear Algebra, Data Structures and Algorithms	

PROJECTS

AI Resume Match & Optimizer <i>Python, NLP, LLMs, FastAPI, pgvector</i>	04/2026 - Present
– Built an end-to-end AI/ML pipeline to parse resumes and job descriptions, extract structured features, and compute match quality with semantic similarity and lexical overlap.	
– Implemented sentence-transformer embeddings and pgvector-based retrieval to improve skill matching, ATS keyword coverage analysis, and requirement-level evidence generation.	
– Developed truth-guarded resume optimization to validate AI-generated rewrite suggestions and reduce unsupported or hallucinated claims.	
– Added evaluation workflows for parsing, matching, semantic retrieval, and guardrail performance to make the system more reliable and production-oriented.	
Personal AI Second Brain <i>Python, FastAPI, LangGraph, ChromaDB, PostgreSQL</i>	02/2026 - Present
– Built hybrid RAG pipeline with BM25 and dense vector search fused via RRF and reranked with Jina cross-encoder.	
– Designed LangGraph agent with LLM-graded self-correction loops to reduce hallucination and enforce citation accuracy.	
– Engineered personal memory layer with entity extraction, time-decay salience scoring and knowledge graph context injection.	
– Evaluated system with custom RAG metrics on a 40+ QA dataset using MLflow-style experiment tracker and cost ledger.	

TECHNICAL SKILLS

Programming Languages: Python, SQL, C++, JavaScript, Java
Machine Learning: Regression, Classification, Clustering, Feature Engineering, Model Evaluation
Deep Learning: Neural Networks, CNN, RNN, LSTM, Transfer Learning, Transformers
Libraries/Frameworks: NumPy, Pandas, Scikit-learn, TensorFlow, PyTorch, OpenCV, Matplotlib, Seaborn
Tools: Git, GitHub, Jupyter Notebook, Google Colab, Docker, Linux
Databases: MySQL, PostgreSQL, MongoDB, ChromaDB, Qdrant

CERTIFICATIONS

Machine Learning Specialization <i>Coursera / Stanford Online</i>	2025
Deep Learning Specialization <i>Coursera / DeepLearning.AI</i>	2025
Python for Data Science <i>Kaggle Learn / DataCamp</i>	2024

LANGUAGES

English: Intermediate / Upper-intermediate; able to read technical documentation and research papers
